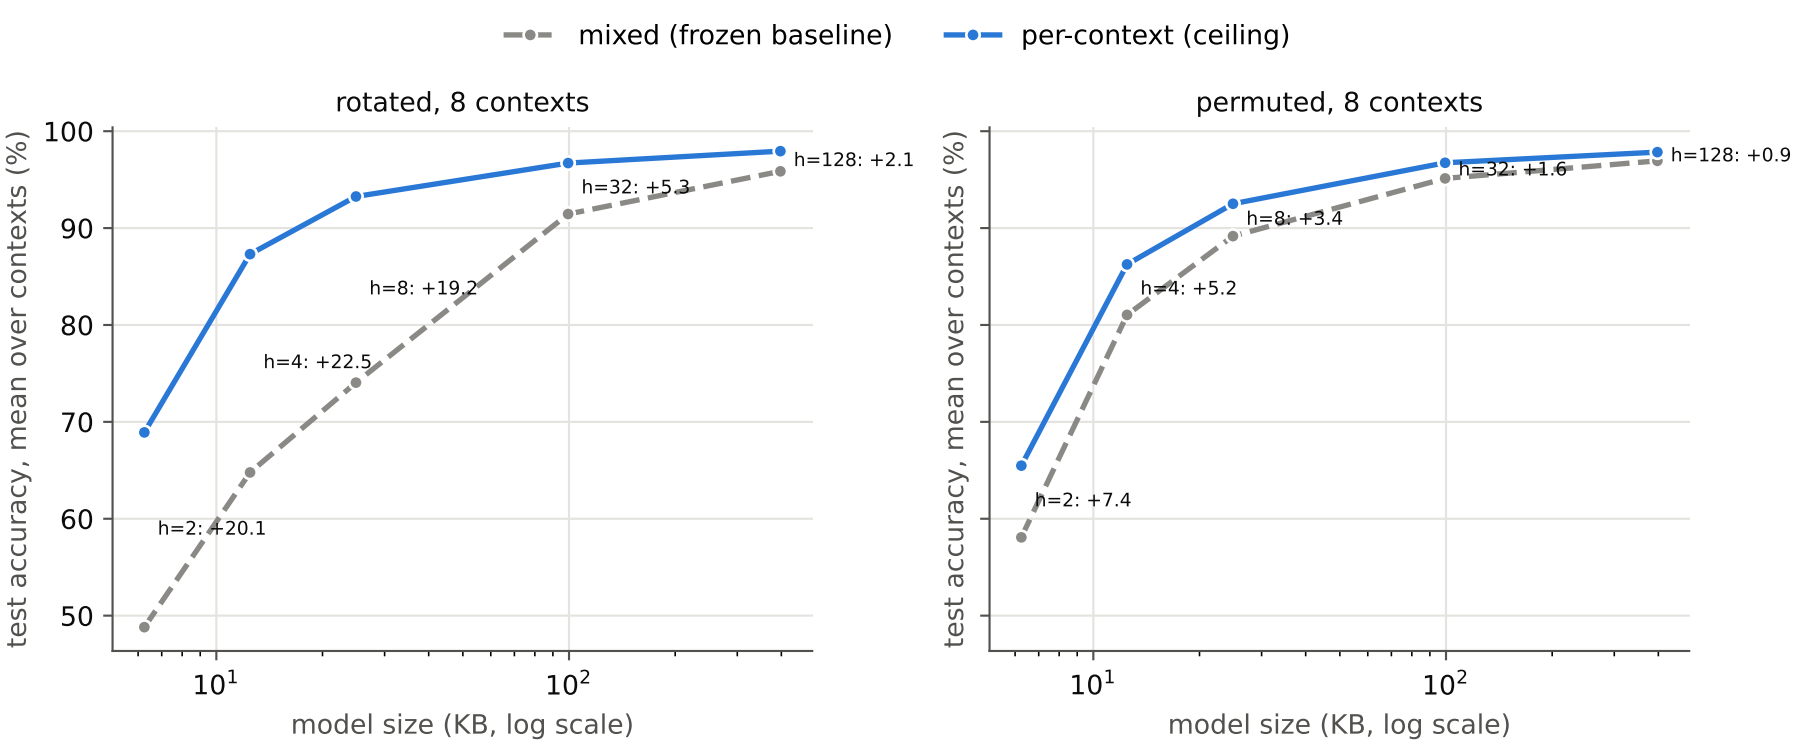
mixed (frozen baseline)
per-context (ceiling)
rotated, 8 contexts
permuted, 8 contexts
test accuracy, mean over contexts (%)
model size (KB, log scale)
h=2: +20.1
h=4: +22.5
h=8: +19.2
h=32: +5.3
h=128: +2.1
h=2: +7.4
h=4: +5.2
h=8: +3.4
h=32: +1.6
h=128: +0.9