Research Highlights

- Dynamic evaluation turns a small model's weights into a cache of recent context
- This memory needs input locality and is tied mostly to the book being read
- Other text evicts it: the memory halves within a few thousand tokens
- The memory sits mostly in the MLPs and costs out-of-domain knowledge
- On long books, a 14M model closes half the gap to a frozen model twice its size